

# Final Project

For the following project, you will be working with a movie dataset. The dataset is [here](#). The dataset columns are as follows:

- Title: The movie's title
- Genre: The movie's genre
- Stars: The number of famous actors in the movie
- Runtime: The length of the movie's runtime
- Budget: How much was spent on filming the movie (in millions)
- Promo: How much money was spent promoting the movie (in millions)
- Season: The season in which the movie was released
- Rating: The movie's rating
- R1: Reviewer 1's review
- R1: Reviewer 2's review
- R1: Reviewer 3's review

And the target variable:

- Success: Whether the film was a success or a flop

Fill in the answers to questions in the text field, and show your code below.

## Data loading

Load the data

```
import pandas as pd

df = pd.read_csv('CMSC320FinalProjectData.csv')

movie_counts = df['Title'].value_counts()
print(movie_counts)
print(df.dtypes)
```

"Whispers of Redemption"	4
"Shadow Strike"	4
"Fading Memories"	4
"Aurora's Legacy"	4
"Echoes of Tomorrow"	4
..	
"Final Strike"	1
"Thunderbolt Fury"	1
"Code of Honor: Dark Redemption"	1
"Warrior's Vow"	1

```
"Thunderstorm Showdown"          1
Name: Title, Length: 460, dtype: int64
Unnamed: 0          int64
Title              object
Runtime           int64
Stars            int64
Year            int64
Budget          float64
Promo          float64
Season          object
Rating          object
Genre          object
R1             object
R2             object
R3             object
Success         bool
dtype: object
```

## Data Cleaning

List the three biggest data errors below, with a summary of how you fixed them and why you choose that method:

- The first biggest data error I saw was that all 3 reviews for every movie is a string and it is hard to do data science stuff on these strings. My first instinct was to read each string and give it a numerical rating myself, but I quickly realized that it would take way too long. Then I decided (by reading through Piazza) it is a lot easier and a lot cooler to import a sentiment analyzer and have it search for some "red flag" words that determine if the reviewer felt negatively towards or positively towards the film.
- The second biggest data error I saw was that there were some clear outliers in the Runtime and Stars column. Specifically, there were some movies with a 0 minute runtime. To alleviate this, I replaced all the 0 minute runtimes with the average runtime of all the movies. Additionally, the Stars column had some movies that had 100 stars which isn't possible. To alleviate this, I replaced the 100 stars with the number of stars for a movie with that same production budget.
- The last biggest data error I saw was that the data needs to account for inflation. This is a necessary step to truly understand the data when considering the production/promotional budgets of the films and seeing if that has any connection with the success of the films.

First take care of data error 1 (the reviews not being usable since right now they are just long strings)

```
!pip install nltk
import nltk
nltk.download('vader_lexicon')
```

```
Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-packages (3.8.1)
Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from nltk) (8.1.7)
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from nltk) (1.3.2)
Requirement already satisfied: regex<=2021.8.3 in /usr/local/lib/python3.10/dist-packages (from nltk) (2023.6.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from nltk) (4.66.1)
```

```
[nltk_data] Downloading package vader_lexicon to /root/nltk_data...
[nltk_data]   Package vader_lexicon is already up-to-date!
```

True

```
from nltk.sentiment.vader import SentimentIntensityAnalyzer
from sklearn.preprocessing import LabelEncoder

def analyze_review(text, pos_threshold=0.1, neg_threshold=-0.1):
    analyzer = SentimentIntensityAnalyzer()
    sentiment_scores = analyzer.polarity_scores(text)

    if sentiment_scores['compound'] >= pos_threshold:
        return "Positive"
    else:
        return "Negative"

# example usage to make sure this sentiment analyzer works
review = "I thought this movie wasn't good "
sentiment = analyze_review(review)
print(f"Sentiment: {sentiment}")

# Make a for loop to loop through the 3 columns of reviews and run the sentiment analyzer.
# First I am going to make new columns and verify the first few entries to make sure
# the sentiment analyzer is working properly the way I want it to.
columns_to_analyze = ['R1', 'R2', 'R3']
for col in columns_to_analyze:
    df[f'Sentiment_{col}'] = df[col].apply(analyze_review)

print(df)

# display the counts of positive and negative for each column just for me to see, this doesn't impact my project
for col in columns_to_analyze:
    positive_count = (df[f'Sentiment_{col}'] == 'Positive').sum()
    negative_count = (df[f'Sentiment_{col}'] == 'Negative').sum()
    positive_counts = positive_count
    negative_counts = negative_count
```

```

for col in columns_to_analyze:
    print(f"Column '{col}':")
    print(f"Positive Reviews: {positive_counts}")
    print(f"Negative Reviews: {negative_counts}")
    print()

# then I encode the data so it will be usable for me. 1 for positive
# reviews and 0 for negative reviews.

# Create a LabelEncoder instance
label_encoder = LabelEncoder()

# Assuming you have already created your DataFrame 'df' with the
# 'Sentiment' columns

# Define the columns you want to encode
columns_to_encode = ['Sentiment_R1', 'Sentiment_R2', 'Sentiment_R3']

# Apply label encoding to the specified columns
for col in columns_to_encode:
    df[col] = label_encoder.fit_transform(df[col])

# Display the updated DataFrame
print(df)

```

```

Sentiment: Negative

```

	Unnamed: 0	Title	Runtime	Stars	Year	\
0	0	"Love in the Inbox"	126	1	2020	
1	1	"Coffee Shop Serendipity"	131	0	2020	
2	2	"The Wedding Date Dilemma"	132	4	2000	
3	3	"Heartstrings and Highways"	132	1	2015	
4	4	"Falling for Cupid"	119	1	2015	
...	...	...	...	...	...	
535	535	"Shadow Strike"	128	3	2021	
536	536	"Riot Protocol"	123	1	2018	
537	537	"Deadlock Vendetta"	121	1	2003	
538	538	"Blade Runner Protocol"	124	1	2007	
539	539	"Eagle Eye Blitz"	126	0	2022	

	Budget	Promo	Season	Rating	Genre	\
0	6.679387e+07	73.543754	Winter	PG	Romantic Comedy	
1	4.667863e+01	33.572003	Fall	PG	Romantic Comedy	
2	3.639134e+01	54.561523	Summer	PG	Romantic Comedy	
3	9.324732e+01	59.714535	Winter	PG13	Romantic Comedy	
4	9.213021e+01	67.643810	Fall	PG13	Romantic Comedy	
...	...	...	...	...	...	
535	6.489702e+01	91.445593	Fall	PG	Action	
536	3.098935e+01	46.045408	Summer	R	Action	
537	4.857255e+01	63.660912	Summer	PG	Action	

538	1.364682e+02	188.513344	Summer	R	Action
539	1.276156e+02	158.496055	Summer	PG	Action
R1 \					
0	"An unconvincing portrayal of suspense that fa...				
1	"A movie that feels disjointed and fails to co...				
2	"An underwhelming cinematic effort with unconv...				
3	"A film that fails to resonate due to its lack...				
4	"A movie that struggles to evoke any genuine e...				
..	...				
535	"Weak and contrived dialogue that lacks authen...				
536	"A film that lingers in the memory."				
537	"A lack of cohesion in the storytelling that m...				
538	"An evocative journey that captivates the soul."				
539	"A movie that speaks a universal language."				
R2 \					
0	"An uninspired plotline that lacks coherence a...				
1	"An attempt at humor that lacks cleverness and...				
2	"An emotionally resonant movie that connects u...				
3	"A beautifully crafted narrative that unfolds ...				
4	"A testament to the power of storytelling, lea...				
..	...				
535	"An overemphasis on spectacle over substance t...				
536	"A celebration of life and its intricacies."				
537	"An overly convoluted plot that confuses rathe...				
538	"An absence of emotional depth, resulting in a...				
539	"Uninspired storytelling that fails to ignite ...				
R3 Success					
Sentiment_R1 \					
0	"A visually captivating masterpiece that mesme...				False
Negative					
1	"A timeless classic that continues to enchant ...				False
Negative					
2	"A cinematic triumph that surpasses boundaries...				False
Negative					
3	"An uninspired portrayal of drama that feels s...				False
Negative					
4	"An uplifting film that leaves a profound impa...				False
Negative					
..	...				...
...					
535	"Unremarkable cinematography that fails to cre...				False
Negative					
536	"A cinematic tour de force that enchants."				True
Negative					
537	"Flat and unconvincing performances that fail ...				False
Negative					
538	"A triumph in storytelling and authenticity."				True

Negative  
 539 "A visually striking and emotionally rich film." True  
 Negative

	Sentiment_R2	Sentiment_R3
0	Negative	Positive
1	Positive	Negative
2	Negative	Positive
3	Positive	Negative
4	Positive	Negative
..	...	...
535	Negative	Negative
536	Negative	Negative
537	Negative	Negative
538	Positive	Positive
539	Negative	Positive

[540 rows x 17 columns]

Column 'R1':

Positive Reviews: 204

Negative Reviews: 336

Column 'R2':

Positive Reviews: 204

Negative Reviews: 336

Column 'R3':

Positive Reviews: 204

Negative Reviews: 336

	Unnamed: 0	Title	Runtime	Stars	Year	\
0	0	"Love in the Inbox"	126	1	2020	
1	1	"Coffee Shop Serendipity"	131	0	2020	
2	2	"The Wedding Date Dilemma"	132	4	2000	
3	3	"Heartstrings and Highways"	132	1	2015	
4	4	"Falling for Cupid"	119	1	2015	
..	...		...	...	...	
535	535	"Shadow Strike"	128	3	2021	
536	536	"Riot Protocol"	123	1	2018	
537	537	"Deadlock Vendetta"	121	1	2003	
538	538	"Blade Runner Protocol"	124	1	2007	
539	539	"Eagle Eye Blitz"	126	0	2022	

	Budget	Promo	Season	Rating	Genre	\
0	6.679387e+07	73.543754	Winter	PG	Romantic Comedy	
1	4.667863e+01	33.572003	Fall	PG	Romantic Comedy	
2	3.639134e+01	54.561523	Summer	PG	Romantic Comedy	
3	9.324732e+01	59.714535	Winter	PG13	Romantic Comedy	
4	9.213021e+01	67.643810	Fall	PG13	Romantic Comedy	
..	...	...	...	...	...	

535	6.489702e+01	91.445593	Fall	PG	Action
536	3.098935e+01	46.045408	Summer	R	Action
537	4.857255e+01	63.660912	Summer	PG	Action
538	1.364682e+02	188.513344	Summer	R	Action
539	1.276156e+02	158.496055	Summer	PG	Action

				R1	\
0	"An unconvincing portrayal of suspense that fa...				
1	"A movie that feels disjointed and fails to co...				
2	"An underwhelming cinematic effort with unconv...				
3	"A film that fails to resonate due to its lack...				
4	"A movie that struggles to evoke any genuine e...				
..	...				
535	"Weak and contrived dialogue that lacks authen...				
536	"A film that lingers in the memory."				
537	"A lack of cohesion in the storytelling that m...				
538	"An evocative journey that captivates the soul."				
539	"A movie that speaks a universal language."				

				R2	\
0	"An uninspired plotline that lacks coherence a...				
1	"An attempt at humor that lacks cleverness and...				
2	"An emotionally resonant movie that connects u...				
3	"A beautifully crafted narrative that unfolds ...				
4	"A testament to the power of storytelling, lea...				
..	...				
535	"An overemphasis on spectacle over substance t...				
536	"A celebration of life and its intricacies."				
537	"An overly convoluted plot that confuses rathe...				
538	"An absence of emotional depth, resulting in a...				
539	"Uninspired storytelling that fails to ignite ...				

				R3	Success
Sentiment_R1 \					
0	"A visually captivating masterpiece that mesme...				False
0					
1	"A timeless classic that continues to enchant ...				False
0					
2	"A cinematic triumph that surpasses boundaries...				False
0					
3	"An uninspired portrayal of drama that feels s...				False
0					
4	"An uplifting film that leaves a profound impa...				False
0					
..	...				...
...					
535	"Unremarkable cinematography that fails to cre...				False
0					
536	"A cinematic tour de force that enchants."				True
0					

```

537 "Flat and unconvincing performances that fail ... False
0
538 "A triumph in storytelling and authenticity." True
0
539 "A visually striking and emotionally rich film." True
0

```

	Sentiment_R2	Sentiment_R3
0	0	1
1	1	0
2	0	1
3	1	0
4	1	0
..	...	...
535	0	0
536	0	0
537	0	0
538	1	1
539	0	1

```

[540 rows x 17 columns]

```

Next, I am going to take care of data error 2 (the clear outliers in the Stars and Runtime column)

```

# first calculate the average run time of movies to replace the 0
# minute movie runtimes.
# I don't do anything special for different genres or anything.
average_runtime = df['Runtime'].mean()
print("Average runtime: ", average_runtime)

# then use the lambda function to replace movies with 0 minute runtime
# with the average we calculated
df['Runtime'] = df['Runtime'].apply(lambda x: average_runtime if x ==
0 else x)

# next I am going to replace movies that have 100 stars with the
# number of stars used in a movie with a similar budget.
def replace_100_stars(row, budget_tolerance=10): # set 5 as budget
tolerance because (I think) the budgets are in millions of dollars. I
feel like 5 million is close enough
    if row['Stars'] == 100:
        # Find movies with budgets within the specified range
        (give/take the budget_tolerance)
        similar_budget_movies = df[
            (df['Budget'] >= (row['Budget'] - budget_tolerance)) &
            (df['Budget'] <= (row['Budget'] + budget_tolerance)) &
            (df['Stars'] != 100)
        ]
        if not similar_budget_movies.empty:

```



```

        return similar_budget_movies['Stars'].mean()
    else:
        # If no similar-budget movie is found, replace with the
        # average number of stars for all movies
        return df['Stars'].mean()

    return row['Stars']

df['Stars'] = df.apply(replace_100_stars, axis=1)

# Display the updated DataFrame
print(df)

```

Average runtime: 130.05835390946504

	Unnamed: 0	Title	Runtime	Stars	Year	\
0	0	"Love in the Inbox"	126.0	1.0	2020	
1	1	"Coffee Shop Serendipity"	131.0	0.0	2020	
2	2	"The Wedding Date Dilemma"	132.0	4.0	2000	
3	3	"Heartstrings and Highways"	132.0	1.0	2015	
4	4	"Falling for Cupid"	119.0	1.0	2015	
...	...	...	...	...	...	
535	535	"Shadow Strike"	128.0	3.0	2021	
536	536	"Riot Protocol"	123.0	1.0	2018	
537	537	"Deadlock Vendetta"	121.0	1.0	2003	
538	538	"Blade Runner Protocol"	124.0	1.0	2007	
539	539	"Eagle Eye Blitz"	126.0	0.0	2022	

	Budget	Promo	Season	Rating	Genre	\
0	6.679387e+07	73.543754	Winter	PG	Romantic Comedy	
1	4.667863e+01	33.572003	Fall	PG	Romantic Comedy	
2	3.639134e+01	54.561523	Summer	PG	Romantic Comedy	
3	9.324732e+01	59.714535	Winter	PG13	Romantic Comedy	
4	9.213021e+01	67.643810	Fall	PG13	Romantic Comedy	
...	...	...	...	...	...	
535	6.489702e+01	91.445593	Fall	PG	Action	
536	3.098935e+01	46.045408	Summer	R	Action	
537	4.857255e+01	63.660912	Summer	PG	Action	
538	1.364682e+02	188.513344	Summer	R	Action	
539	1.276156e+02	158.496055	Summer	PG	Action	

	R1	\
0	"An unconvincing portrayal of suspense that fa...	
1	"A movie that feels disjointed and fails to co...	
2	"An underwhelming cinematic effort with unconv...	
3	"A film that fails to resonate due to its lack...	
4	"A movie that struggles to evoke any genuine e...	
...	...	
535	"Weak and contrived dialogue that lacks authen...	
536	"A film that lingers in the memory."	
537	"A lack of cohesion in the storytelling that m...	

538 "An evocative journey that captivates the soul."  
539 "A movie that speaks a universal language."

R2 \

0 "An uninspired plotline that lacks coherence a...  
1 "An attempt at humor that lacks cleverness and...  
2 "An emotionally resonant movie that connects u...  
3 "A beautifully crafted narrative that unfolds ...  
4 "A testament to the power of storytelling, lea...  
.. ...  
535 "An overemphasis on spectacle over substance t...  
536 "A celebration of life and its intricacies."  
537 "An overly convoluted plot that confuses rathe...  
538 "An absence of emotional depth, resulting in a...  
539 "Uninspired storytelling that fails to ignite ...

R3 Success

Sentiment\_R1 \

0 "A visually captivating masterpiece that mesme... False  
0  
1 "A timeless classic that continues to enchant ... False  
0  
2 "A cinematic triumph that surpasses boundaries... False  
0  
3 "An uninspired portrayal of drama that feels s... False  
0  
4 "An uplifting film that leaves a profound impa... False  
0  
.. ...  
...  
535 "Unremarkable cinematography that fails to cre... False  
0  
536 "A cinematic tour de force that enchants." True  
0  
537 "Flat and unconvincing performances that fail ... False  
0  
538 "A triumph in storytelling and authenticity." True  
0  
539 "A visually striking and emotionally rich film." True  
0

	Sentiment_R2	Sentiment_R3
0	0	1
1	1	0
2	0	1
3	1	0
4	1	0
..	...	...
535	0	0
536	0	0

537	0	0
538	1	1
539	0	1

[540 rows x 17 columns]

Next, I will tackle data error 3 (accounting for inflation)

```
import numpy as np

# first print the value counts of each year in the dataframe so I know
which years i have to account for inflation.
year_counts = df['Year'].value_counts().reset_index()
year_counts.columns = ['Year', 'Count']
sorted_movies_df = year_counts.sort_values(by='Year')

print(sorted_movies_df)
# now I can see that I have to adjust for inflation for years 2000-
2023

# I get my factors of inflation from this website:
https://www.in2013dollars.com/us/inflation/2003?amount=1
inflation_factors = {
    # I'm sure there is a much more elegant way to do this but I don't
    know much about inflation curves so to be save I will do it like this
    2000: 1.79,
    2001: 1.74,
    2002: 1.71,
    2003: 1.67,
    2004: 1.63,
    2005: 1.58,
    2006: 1.53,
    2007: 1.48,
    2008: 1.43,
    2009: 1.43,
    2010: 1.41,
    2011: 1.37,
    2012: 1.34,
    2013: 1.32,
    2014: 1.30,
    2015: 1.30,
    2016: 1.28,
    2017: 1.26,
    2018: 1.22,
    2019: 1.20,
    2020: 1.19,
    2021: 1.14,
    2022: 1.05,
    2023: 1.0, # No inflation adjustment needed for 2023
```

```

}

# Define the range of years you want to adjust for
years_to_adjust = range(2000, 2024)

# tterate through each year and adjust 'Budget' and 'Promo' columns
depending on the year
for year in years_to_adjust:
    if year in inflation_factors:
        inflation_factor = inflation_factors[year]
        # update budget column
        df.loc[df['Year'] == year, 'Budget'] *= inflation_factor
        # update promo column
        df.loc[df['Year'] == year, 'Promo'] *= inflation_factor

# I noticed the budget column has some crazy outliers to I am going to
transform it to reduce the impact of the super high values
df['Budget'] = np.log1p(df['Budget'])

# check data frame
print(df)

```

	Year	Count
19	2000	17
18	2001	18
0	2002	37
13	2003	20
9	2004	22
22	2005	15
11	2006	22
3	2007	30
4	2008	30
17	2009	18
12	2010	20
15	2011	19
6	2012	26
23	2013	13
7	2014	26
5	2015	27
14	2016	20
21	2017	16
8	2018	22
2	2019	30
10	2020	22
16	2021	19
1	2022	34
20	2023	17

	Unnamed: 0	Title	Runtime	Stars	Year
Budget	\				

0	0	"Love in the Inbox"	126.0	1.0	2020
18.365028					
1	1	"Coffee Shop Serendipity"	131.0	0.0	2020
4.206208					
2	2	"The Wedding Date Dilemma"	132.0	4.0	2000
4.767302					
3	3	"Heartstrings and Highways"	132.0	1.0	2015
5.066309					
4	4	"Falling for Cupid"	119.0	1.0	2015
5.054333					
..	...	...	...	...	...
...					
535	535	"Shadow Strike"	128.0	3.0	2021
4.446645					
536	536	"Riot Protocol"	123.0	1.0	2018
3.852794					
537	537	"Deadlock Vendetta"	121.0	1.0	2003
4.916061					
538	538	"Blade Runner Protocol"	124.0	1.0	2007
5.703515					
539	539	"Eagle Eye Blitz"	126.0	0.0	2022
4.953686					
	Promo	Season	Rating	Genre	\
0	104.145310	Winter	PG	Romantic Comedy	
1	47.541313	Fall	PG	Romantic Comedy	
2	174.820577	Summer	PG	Romantic Comedy	
3	100.917564	Winter	PG13	Romantic Comedy	
4	114.318039	Fall	PG13	Romantic Comedy	
..	...	...	...	...	
535	118.842692	Fall	PG	Action	
536	68.533985	Summer	R	Action	
537	177.543919	Summer	PG	Action	
538	412.919629	Summer	R	Action	
539	174.741900	Summer	PG	Action	
					R1 \
0	"An unconvincing portrayal of suspense that fa..."				
1	"A movie that feels disjointed and fails to co..."				
2	"An underwhelming cinematic effort with unconv..."				
3	"A film that fails to resonate due to its lack..."				
4	"A movie that struggles to evoke any genuine e..."				
..	...				
535	"Weak and contrived dialogue that lacks authen..."				
536	"A film that lingers in the memory."				
537	"A lack of cohesion in the storytelling that m..."				
538	"An evocative journey that captivates the soul."				
539	"A movie that speaks a universal language."				
					R2 \

```

0 "An uninspired plotline that lacks coherence a...
1 "An attempt at humor that lacks cleverness and...
2 "An emotionally resonant movie that connects u...
3 "A beautifully crafted narrative that unfolds ...
4 "A testament to the power of storytelling, lea...
..
535 "An overemphasis on spectacle over substance t...
536 "A celebration of life and its intricacies."
537 "An overly convoluted plot that confuses rathe...
538 "An absence of emotional depth, resulting in a...
539 "Uninspired storytelling that fails to ignite ...

```

R3 Success

Sentiment\_R1 \

```

0 "A visually captivating masterpiece that mesme... False
0
1 "A timeless classic that continues to enchant ... False
0
2 "A cinematic triumph that surpasses boundaries... False
0
3 "An uninspired portrayal of drama that feels s... False
0
4 "An uplifting film that leaves a profound impa... False
0
.. ...
...
535 "Unremarkable cinematography that fails to cre... False
0
536 "A cinematic tour de force that enchants." True
0
537 "Flat and unconvincing performances that fail ... False
0
538 "A triumph in storytelling and authenticity." True
0
539 "A visually striking and emotionally rich film." True
0

```

	Sentiment_R2	Sentiment_R3
0	0	1
1	1	0
2	0	1
3	1	0
4	1	0
..	...	...
535	0	0
536	0	0
537	0	0
538	1	1
539	0	1

[540 rows x 17 columns]

## Data Exploration

Does Season have a statistically significant impact on a movie's success?

**p-value:** 0.005716268505111858

```
# we will do a chi squared test since the seasons are different categories
from scipy.stats import chi2_contingency

contingency_table = pd.crosstab(df['Season'], df['Success'])
chi2, p, dof, expected = chi2_contingency(contingency_table)

# i made my significance level .05 because that is commonly chosen significance level
alpha = 0.05

# Interpret the results
if p < alpha:
    print("There is a statistically significant association between season and success.")
    print(p)
else:
    print("There is no statistically significant association between season and success.")
    print(p)

There is a statistically significant association between season and success.
0.005716268505111858
```

Do seasons have a statistically significant difference in their distribution of content ratings?

**p-value:** 0.21507814500508263

```
# we will do another chi squared test since we are still doing seasons
contingency_table = pd.crosstab(df['Season'], df['Rating'])

chi2, p, _, _ = chi2_contingency(contingency_table)
alpha = 0.05
if p < alpha:
    print("There is a statistically significant difference between Season and Rating.")
    print(p)
else:
```

```
print("There is no statistically significant difference between  
Season and Rating.")  
print(p)
```

```
There is no statistically significant difference between Season and  
Rating.  
0.21507814500508263
```

Who is the harshest critic (highest percent of negative reviews)?

**Critic:** Reviewer one is the harshest critic. They have the highest percent of negative reviews.

```
# keep count of number of negative reviews (according to sentiment  
analyzer)  
R1_count = 0  
R2_count = 0  
R3_count = 0  
  
max_negative_count = 0  
total_reviews = len(df)  
  
# iterate through each row of the dataframe and increment counts when  
appropriate  
for index, row in df.iterrows():  
    # Count negative reviews (0 values) for each reviewer  
    R1_count += row['Sentiment_R1'] == 0  
    R2_count += row['Sentiment_R2'] == 0  
    R3_count += row['Sentiment_R3'] == 0  
  
    # Calculate the total count of negative reviews for the current  
    critic  
    total_negative_count = R1_count + R2_count + R3_count  
  
    # Check if the current critic has more negative reviews than the  
    previous maximum  
    if total_negative_count > max_negative_count:  
        max_negative_count = total_negative_count  
        harshest_critic = index # Use the index (row number) as a  
        pseudo-critic name  
  
# Print the harshest critic and their counts for each reviewer  
print("Percentage of negative reviews for each reviewer:")  
print("R1:", (R1_count/total_reviews) * 100, "%")  
print("R1:", (R2_count/total_reviews) * 100, "%")  
print("R1:", (R3_count/total_reviews) * 100, "%")  
  
Percentage of negative reviews for each reviewer:  
R1: 79.81481481481481 %  
R1: 69.62962962962963 %  
R1: 62.22222222222222 %
```



What is the covariance between promotional budget and the filming budget?

**Cov:** 56.423789973820185

```
covariance = df['Promo'].cov(df['Budget'])
print("Covariance of promo and budget columns: ", covariance)

Covariance of promo and budget columns: 56.423789973820185
```

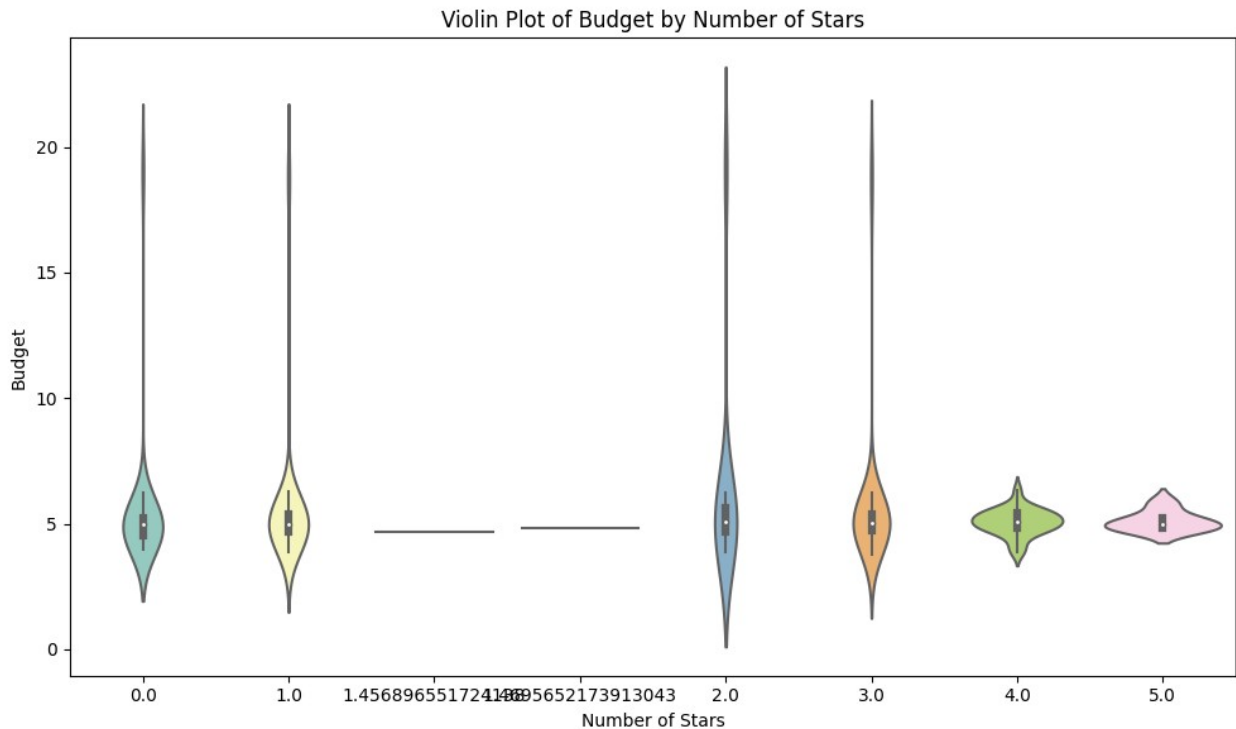
## Data Visualization

Create a chart that compares the distribution of the budget for each different number of stars. (It does not need to be particularly appealing.)

```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
sns.violinplot(data=df, x='Stars', y='Budget', palette='Set3')
plt.title('Violin Plot of Budget by Number of Stars')
plt.xlabel('Number of Stars')
plt.ylabel('Budget')
plt.tight_layout()
plt.show()


# this isn't a very pretty graph I know. Also please remember that I
# transformed the budget column!
# so this graph more so gives us a general relationship idea of
# budget/stars
```



Create a graph showing the average movie budget over time.

```
average_budget_by_year = df.groupby('Year')['Budget'].mean()

# Create a line plot to show the average movie budget over time
plt.figure(figsize=(12, 6)) # Adjust the figure size as needed
plt.plot(average_budget_by_year.index, average_budget_by_year.values,
marker='o', linestyle='-', color='g')
plt.title('Average Movie Budget Over Time')
plt.xlabel('Year')
plt.ylabel('Average Budget (Millions of Dollars)')
plt.grid(True)
plt.tight_layout()
plt.show()
```



## Feature Engineering

List any features you choose to create (if you are creating many features based on one column, you do not need to list them separately.) You are not required to create any features if you do not wish to. You may create any number of additional features.

I made lots of new columns and features because I was very afraid to modify the original data.

- I made new columns (EX: Sentiment\_R1) for each reviewer where I encoded the sentiment I got, so reviews that were analyzed to be Negative were represented by a 0 and those that were analyzed to be Positive were represented by a 1.
- Another new column I made was AverageSentiment that took the average of the 3 new columns I made in the previous bullet point which made it easier to see what the overall sentiment towards the movie.
- There were also many columns I edited (runtime, stars, etc)
- I also made the year counts columns to see how many movies were released in each year and see what years were included in the data frame so I knew which years I had to get the inflation factor for.

For Modeling/Testing, I won't use the R1, R2, and R3 columns since I have the sentiment versions with integers for that. So I don't have to do any encoding.

Since we know seasons don't have a big impact on success, I won't include that either.

Since there are many categories for rating and genre, I won't include them because one hot encoding isn't good for a column with many categories.

# Modeling

Create a model of your choice.

**Model type choosen:** Random forest

```
print(df.dtypes)

# going to make a copy of the dataframe with the columns I want to
include for the testing and stuff
selected_columns = ['Title',
                    'Runtime',
                    'Stars',
                    'Year',
                    'Budget',
                    'Promo',
                    'Success',
                    'Sentiment_R1',
                    'Sentiment_R2',
                    'Sentiment_R3',
                    'AverageSentiment']

selected_df = df[selected_columns].copy()
print(selected_df.head())
```

```
Unnamed: 0      int64
Title           object
Runtime        float64
Stars          float64
Year           int64
Budget         float64
Promo          float64
Season         object
Rating         object
Genre          object
R1             object
R2             object
R3             object
Success        bool
Sentiment_R1   int64
Sentiment_R2   int64
Sentiment_R3   int64
AverageSentiment float64
dtype: object
```

	Title	Runtime	Stars	Year	Budget
Promo \					
0	"Love in the Inbox"	126.0	1.0	2020	18.365028
104.145310					
1	"Coffee Shop Serendipity"	131.0	0.0	2020	4.206208

```

47.541313
2  "The Wedding Date Dilemma"    132.0    4.0  2000    4.767302
174.820577
3  "Heartstrings and Highways"    132.0    1.0  2015    5.066309
100.917564
4      "Falling for Cupid"    119.0    1.0  2015    5.054333
114.318039

```

	Success	Sentiment_R1	Sentiment_R2	Sentiment_R3	AverageSentiment
0	False	0	0	1	0.333333
1	False	0	1	0	0.333333
2	False	0	0	1	0.333333
3	False	0	1	0	0.333333
4	False	0	1	0	0.333333

## Testing

Shuffle your data and break it into a 10% test set and 90% training set. Show your model's accuracy on the test set. In order to get full credit, the model's accuracy must be higher than 50%.

**Model accuracy:** Before adjusting threshold in next step: 0.8148148148148148

After adjusting the threshold in the next step (accuracy is still > 50%):  
0.7962962962962963

```

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

X = selected_df.drop(columns=['Success', 'Title']) # all the columns
to be features (everything except for success)
y = selected_df['Success'] # target !!

# splitting. I think this shuffles for me according to piazza.
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.10, random_state=40)

# make my random forest classifier
clf = RandomForestClassifier(random_state=40)
clf.fit(X_train, y_train)

# makin my predictions

```

```

y_pred = clf.predict(X_test)

# calculate the accuracy of the model on the test set
accuracy = accuracy_score(y_test, y_pred)

# check if the accuracy is higher than 50%
if accuracy > 0.50:
    print(f"Model Accuracy: ", accuracy)
else:
    print(f"Model Accuracy: ", accuracy)

Model Accuracy: 0.8148148148148148

```

Show the confusion matrix for your model. To get full credit, your false positive rate and false negative rate must be under 30%.

(since one of my rates wasn't under 30, i had to change threshold)

**False negative rate:** 0.25

**False positive rate:** 0.19047619047619047

```

from sklearn.metrics import confusion_matrix

# this was the first threshold so like in the last part.
y_pred_default = clf.predict(X_test)

accuracy_default = accuracy_score(y_test, y_pred_default)

confusion_default = confusion_matrix(y_test, y_pred_default)

true_negatives_default, false_positives_default,
false_negatives_default, true_positives_default =
confusion_default.ravel()
false_positive_rate_default = false_positives_default /
(false_positives_default + true_negatives_default)
false_negative_rate_default = false_negatives_default /
(false_negatives_default + true_positives_default)

print("Results BEFORE changing threshold:")
print(f"Accuracy: ", accuracy_default)
print(f"False Positive Rate: ", false_positive_rate_default)
print(f"False Negative Rate: ", false_negative_rate_default)

# lowering the threshold because false negative rate is too high
threshold = 0.3
y_pred_adjusted = (clf.predict_proba(X_test)[:, 1] >
threshold).astype(int)

# do accuracy again
accuracy_adjusted = accuracy_score(y_test, y_pred_adjusted)

```

```

confusion_adjusted = confusion_matrix(y_test, y_pred_adjusted)
true_negatives_adjusted, false_positives_adjusted,
false_negatives_adjusted, true_positives_adjusted =
confusion_adjusted.ravel()
false_positive_rate_adjusted = false_positives_adjusted /
(false_positives_adjusted + true_negatives_adjusted)
false_negative_rate_adjusted = false_negatives_adjusted /
(false_negatives_adjusted + true_positives_adjusted)

print("New results after changing threshold for classification:")
print(f"Accuracy: ", accuracy_adjusted)
print(f"False Positive Rate: ", false_positive_rate_adjusted)
print(f"False Negative Rate: ", false_negative_rate_adjusted)

```

```

Results BEFORE changing threshold:
Accuracy:  0.8148148148148148
False Positive Rate:  0.07142857142857142
False Negative Rate:  0.5833333333333334
New results after changing threshold for classification:
Accuracy:  0.7962962962962963
False Positive Rate:  0.19047619047619047
False Negative Rate:  0.25

```

What was the most important feature for your model? Don't guess, either look up how to check or do your own tests.

Source: [https://scikit-learn.org/stable/auto\\_examples/ensemble/plot\\_forest\\_importances.html](https://scikit-learn.org/stable/auto_examples/ensemble/plot_forest_importances.html)

**Most important feature:** Budget

```

feature_importances = clf.feature_importances_

feature_names = X_train.columns
most_important_idx = feature_importances.argmax()
most_important_feature = feature_names[most_important_idx]
most_important_importance = feature_importances[most_important_idx]

# printing the most important feature and its importance score
print("Most Important Feature: ", most_important_feature)
print("Importance Score: ", most_important_importance)

Most Important Feature:  Budget
Importance Score:  0.211644383883012

```